



LAB 8:

UNSUPERVISED LEARNING AND GANs

University of Washington, Seattle

Spr 2025



OUTLINE

Part 1: Unsupervised Learning

- Supervised vs unsupervised
- Unsupervised learning in with NN

Part 2: Generative Model Taxidermy

- FVBN
- Variational Autoencoder
- GAN

Part 3: Generative Adversarial Networks

- GAN architecture

Part 4: GAN Optimization and Applications

- Competing cost function
- Minmax game optimization
- GAN variations

Part 5: GAN Example

- MNIST generation with Vanilla GAN
- Generator extension with convolution

Part 6: Lab Assignment

- MNIST generation with DCGAN



Unsupervised Learning

Supervised vs Unsupervised

Unsupervised Learning in NN



Supervised vs Unsupervised Learning

Supervised

Data:

{x} x: inputs **WITH** labels

Neural Network Goal:

Minimize specific **error**

Examples: Classification,
Regression, Detection, Prediction



Supervised vs Unsupervised Learning

Supervised

Data:

$\{x\}$ x: inputs **WITH** labels

Neural Network Goal:

Minimize specific **error**


Examples: Classification,
Regression, Detection, Prediction

Regression




What will be the temperature tomorrow?

84°



Fahrenheit



Classification




Will it be hot or cold tomorrow?

COLD

HOT



Fahrenheit





Unsupervised Learning in NN

Unsupervised

Data:

$\{x\}$ x: inputs **WITHOUT** labels

Neural Network Goal:

Learn a **structure** of the data



Unsupervised Learning in NN

Training Data



Training Data $\sim \mathbf{P}_{\text{data}}(\mathbf{x})$



Unsupervised Learning in NN

Training Data



Training Data $\sim \mathbf{P}_{\text{data}}(\mathbf{x})$



Generated Samples

<http://www.whichfaceisreal.com/>



Generate Samples $\sim \mathbf{P}_{\text{model}}(\mathbf{x})$



Unsupervised Learning in NN

Training Data



Training Data $\sim \mathbf{P}_{\text{data}}(\mathbf{x})$

Generated Samples

<http://www.whichfaceisreal.com/>



Generate Samples $\sim \mathbf{P}_{\text{model}}(\mathbf{x})$

Goal: Model estimated density \approx Real world density

Core problem in unsupervised learning



Unsupervised Learning in NN

Unsupervised

Data:

$\{x\}$ x: inputs **WITHOUT** labels

+ No need for labeling → More data

- **Challenge: Cost?**

Neural Network Goal:

Learn a **structure** of the data

+ Has the potential to learn the real world

- **Challenge: Optimization?**




Maximum Likelihood Estimation



Data points

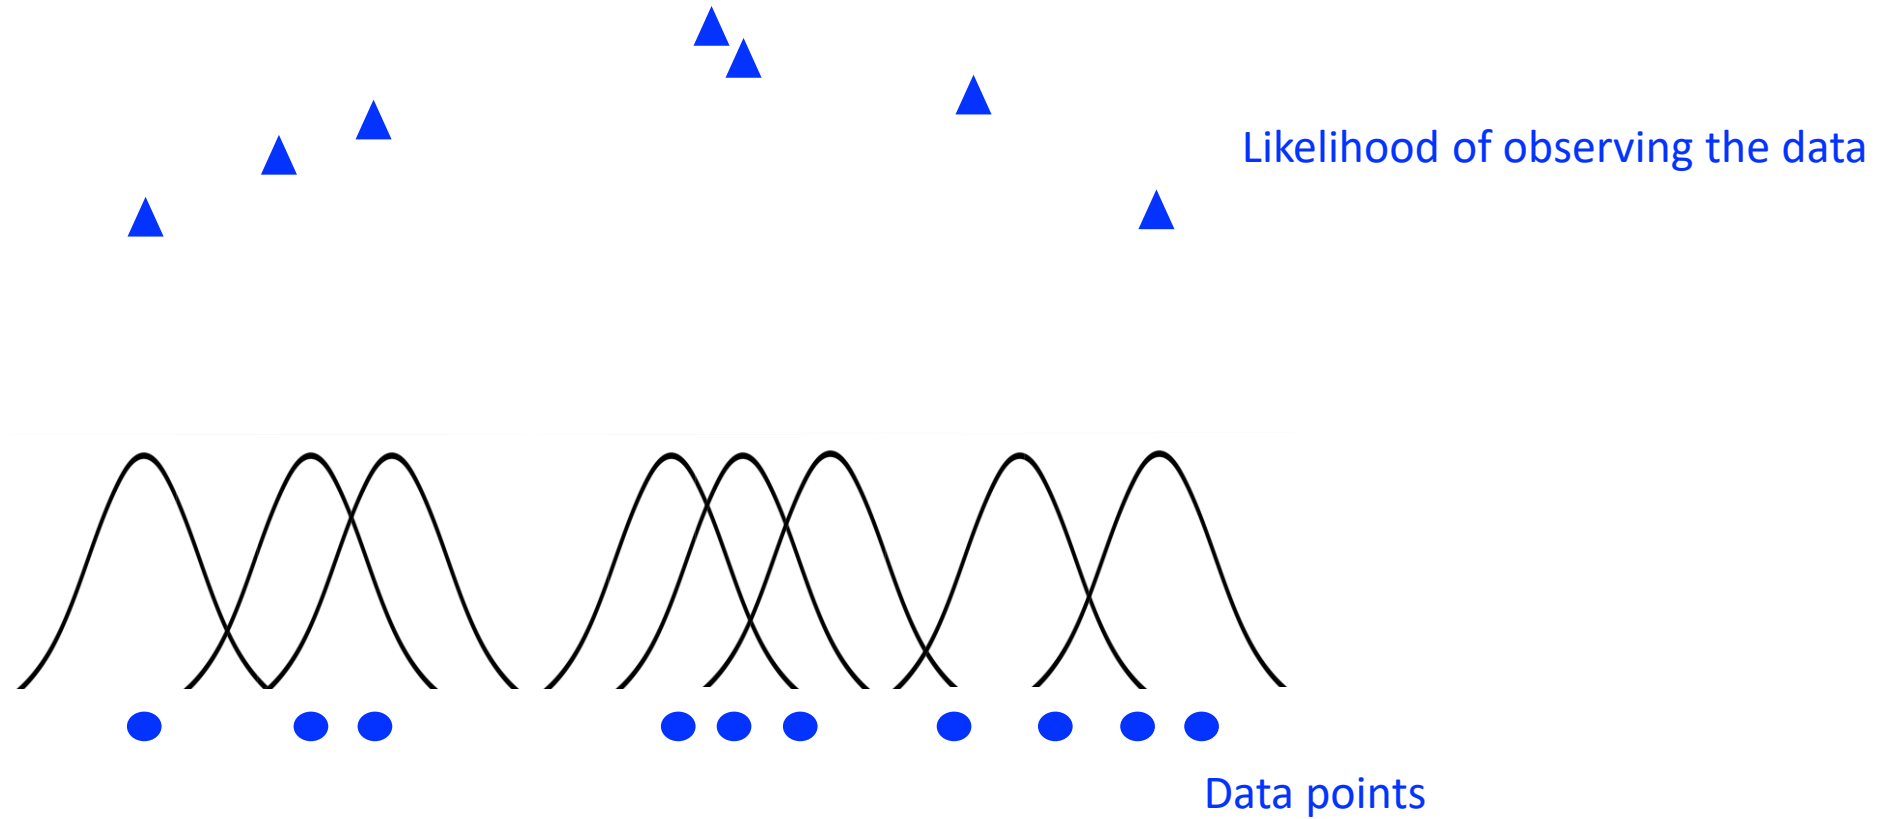


Maximum Likelihood Estimation






Maximum Likelihood Estimation






Unsupervised Learning in NN





Unsupervised Learning in NN




Model params

$$\theta^* = \operatorname{argmax}_{\theta} E_{x \sim p_{data}} \log p_{model}(x|\theta)$$

Goal: Find the optimal distribution $p_{model}(x|\theta)$ that best fit the data



Unsupervised Learning in NN



Model params

$$\theta^* = \underset{\theta}{\operatorname{argmax}} E_{x \sim p_{\text{data}}} \log p_{\text{model}}(x|\theta)$$

Explicit – explicitly define and generate P_{model}

Implicit - generate P_{model} without defining P_{model} exactly



Generative Model Taxidermy

Fully Visible BN

Variational Autoencoder

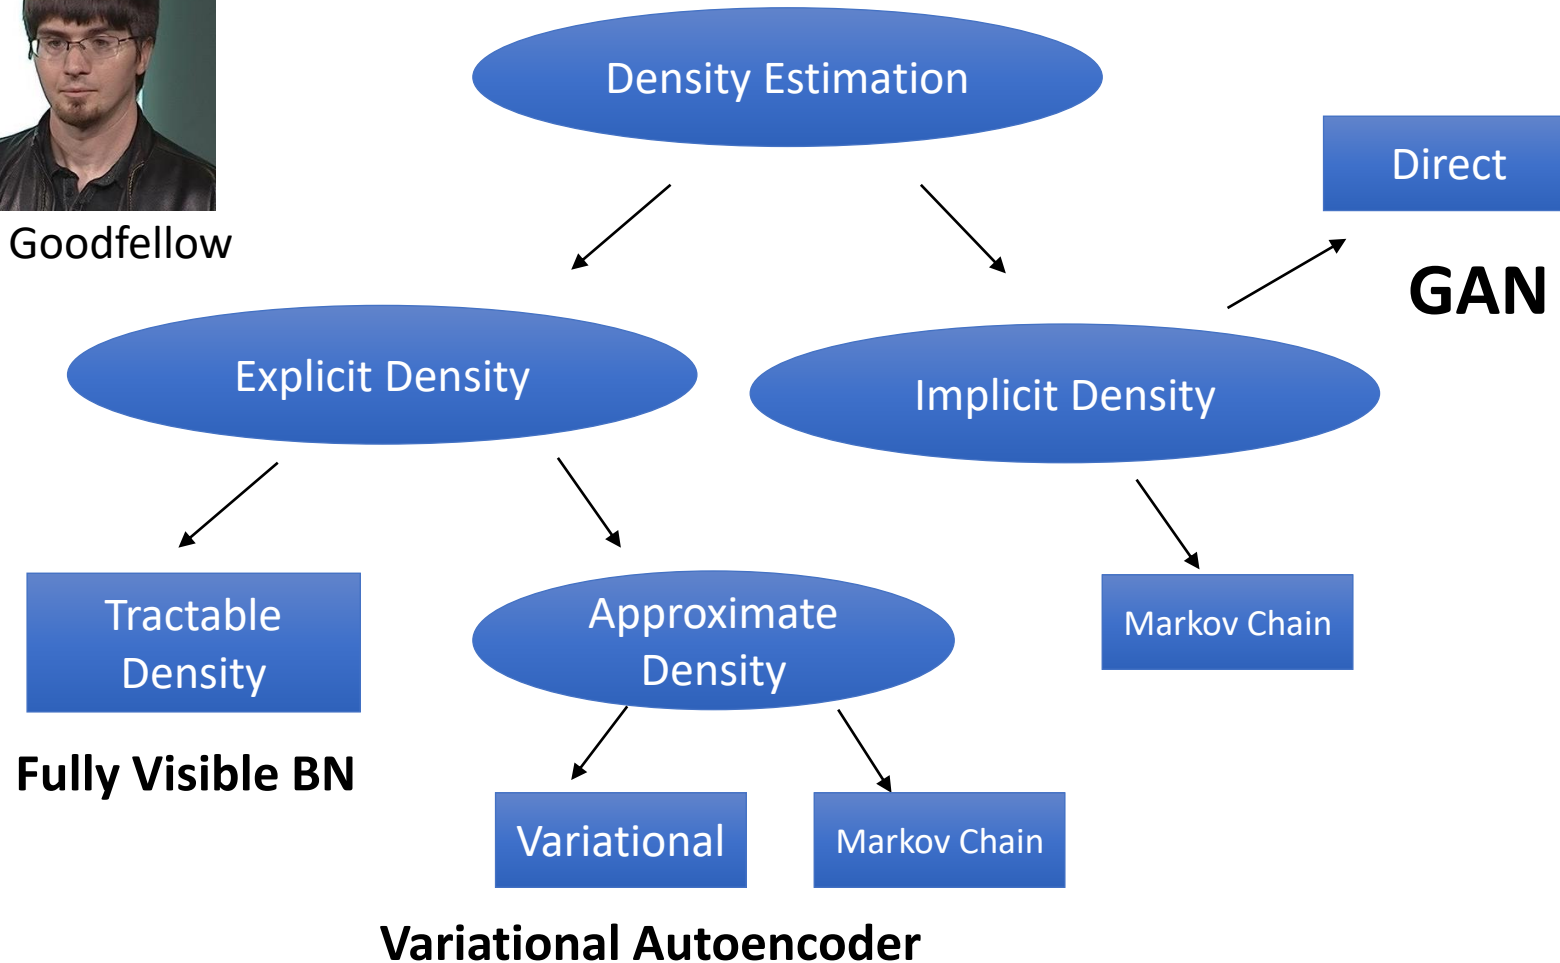
Generative Adversarial Network



Unsupervised Learning in NN



Ian Goodfellow





Fully Visible BN

- Explicitly formula based on chain rule:

$$p_{model}(x) = p_{model}(x_1) \prod_{i=2}^n p_{model}(x_i | x_1, x_2, \dots, x_{i-1})$$

- $O(n)$ generation cost
- No control through hidden variables



Language Model

Language model: probability distribution over sequences of words. Given such a sequence, say of length m , it assigns a probability to the whole sequence.



Language Model

Language model: probability distribution over sequences of words. Given such a sequence, say of length m , it assigns a probability to the whole sequence.

$P(W=\text{NASA will } \mathbf{take} \text{ me to Moon}) = p_1$

$P(W= \text{NASA will } \mathbf{bake} \text{ me to Moon}) = p_2$

$p_2 \ll p_1$



Language Model

Language model: probability distribution over sequences of words. Given such a sequence, say of length m , it assigns a probability to the whole sequence.

$$P(W=\text{NASA will } \mathbf{take} \text{ me to Moon}) = p1$$

$$P(W= \text{NASA will } \mathbf{bake} \text{ me to Moon}) = p2$$

$$p2 \ll p1$$

Chain rule is used to estimate probability:

$$P(w_1 w_2 \dots w_n) = \prod_i P(w_i | w_1 w_2 \dots w_{i-1})$$

$$P(W) = P(\mathbf{NASA}) P(\mathbf{will} | \text{NASA}) P(\mathbf{take} | \text{NASA will}) P(\mathbf{me} | \text{NASA will take}) \\ P(\mathbf{to} | \text{NASA will take me}) P(\mathbf{Moon} | \text{NASA will take me to})$$



Variational Autoencoder

$$p_{model}(\mathbf{x}) = \prod_{i=2}^n p_{model}(x_i | x_1, x_2, \dots, x_{i-1})$$



$$p_{model}(\mathbf{x}) = \int p_{model}(\mathbf{z}) p_{model}(\mathbf{x}|\mathbf{z}) d\mathbf{z}$$



Variational Autoencoder

$$p_{model}(\mathbf{x}) = \prod_{i=2}^n p_{model}(\mathbf{x}_i | \mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_{i-1})$$




$$p_{model}(\mathbf{x}) = \int p_{model}(\mathbf{z}) p_{model}(\mathbf{x} | \mathbf{z}) d\mathbf{z}$$

$p_{model}(\mathbf{x})$ is controlled by hidden state \mathbf{z}




Variational Autoencoder






Variational Autoencoder





GAN



- Instead of sampling from **high dimensional, complex and unknown distribution**
- Sample from **simple distribution**, e.g. normal distribution (random noise) and **find transformation** to the distribution we want to learn.
- **Learn the transformation using a NN**



Generative Adversarial Networks


GAN Architecture

Generator Network

Discriminator Network




GAN





GAN

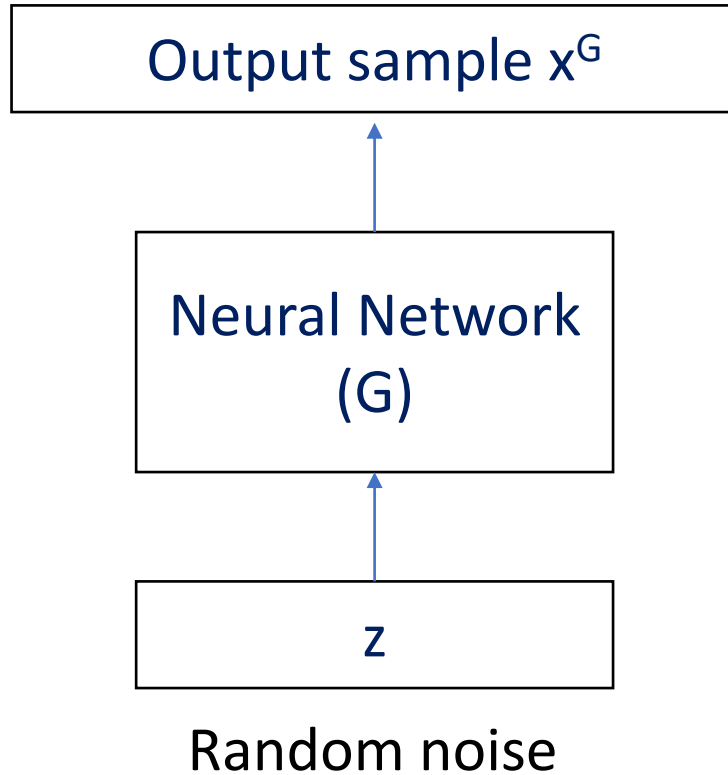


Discriminator – try to distinguish between **x (real)** and **generated (fake)** images

Generator – try to generate **samples and present them as real world** and fool the discriminator



Generator (G)



Training data has distribution \mathbf{p}_{data} .
Sample $\mathbf{x} \sim \mathbf{p}_{\text{data}}$.

Goal: Output sample \mathbf{x}^G is of similar dimensions as \mathbf{x} and distribution \mathbf{p}_{data} .



Examples

Face:



Car:



Bedroom:






Discriminator (D)

Receives input of same dimensions as \mathbf{p}_{data} .

Goal: Distinguish sample from \mathbf{p}_{data} (1) or not (0).





Examples

Discriminator

Face (gen):



0

Car (real):



1

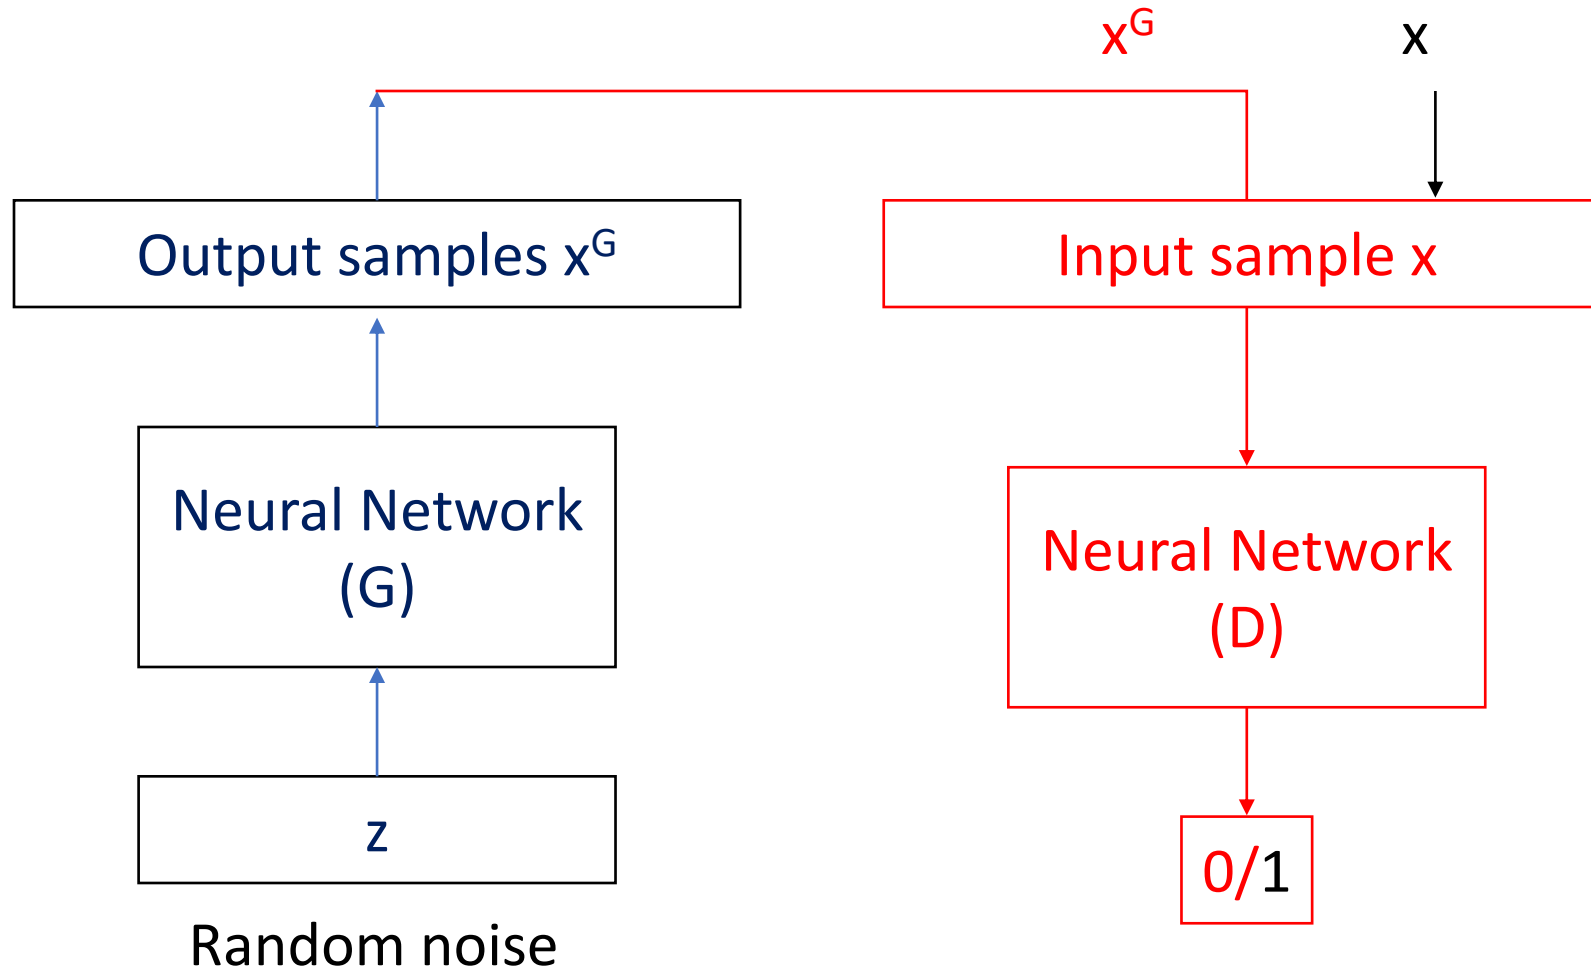
Bedroom (gen):



0



Full Architecture





GAN Optimization and Applications

Competing Loss Functions

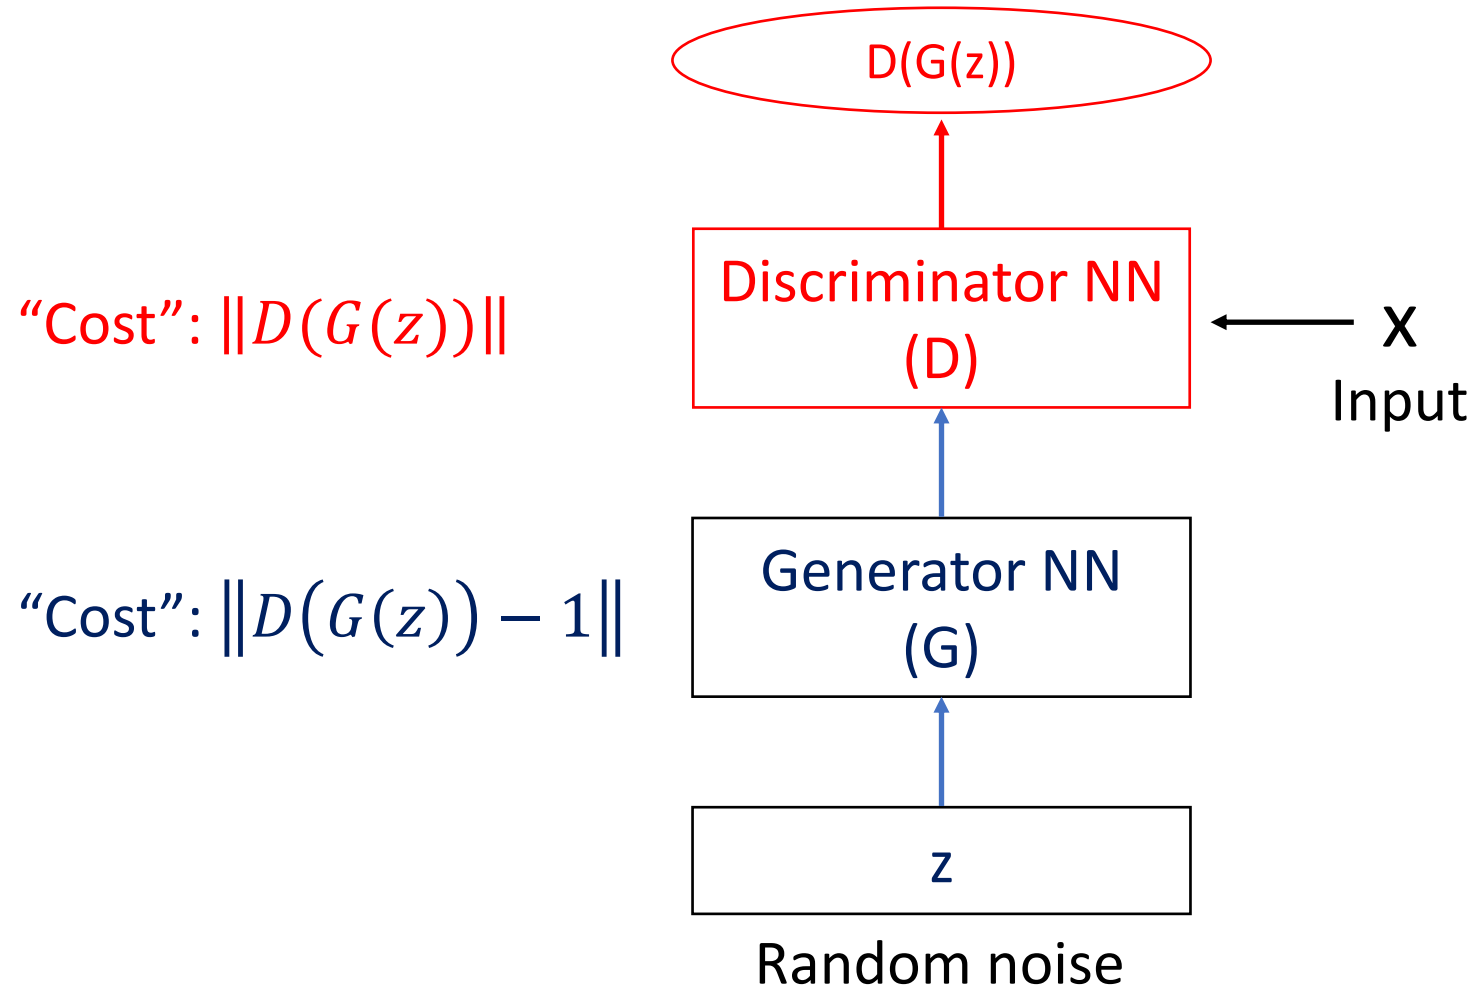
Minmax Game Optimization

Optimization in NN

GAN Applications



Competing cost functions





Competing cost functions

Binary Cross Entropy Loss

$$J^{(D)} = -\frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) - \frac{1}{2} \mathbb{E}_{z \sim p_{model}} \log (1 - D_{\theta_d}(G_{\theta_g}(z)))$$

$$J^{(G)} = -J^{(D)}$$

$$J^{(D)} = -\frac{1}{2} \int p_{data}(x) \log D(x) dx - \frac{1}{2} \int p_{model}(x) \log(1 - D(x)) dx$$



Competing cost functions

Optimal $D(x)$ is


$$D(x) = \frac{p_{data}}{p_{model} + p_{data}}$$

Assumption: p_{model}, p_{data} are nonzero everywhere

Equilibrium: $p_{model} = p_{data}$ then $E(D(x)) = \frac{1}{2}$



Competing cost functions



Discriminator learns an approximation of $p_{\text{data}}(x)/p_{\text{model}}(x)$
vs
learning $p_{\text{model}}(x)$ directly (or indirectly via latent variable models).



Minmax Game Optimization

$$\min_{\theta_g} \max_{\theta_d} \left[\mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p_{model}} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Discriminator output
for real data

Discriminator output
for generated data

Solution:

Saddle point in the parameter space (Nash Equilibrium)

- One player (Discriminator) is at **maximum**,
- Other player (Generator) is at **minimum**



Optimization in NN

- **Gradient ascent** for the discriminator on J

$$J^{(D)} = \frac{1}{2} \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \frac{1}{2} \mathbb{E}_{z \sim p_{model}} \log \left(1 - D_{\theta_d}(G_{\theta_g}(z)) \right)$$

$$\theta_d \leftarrow \arg \min_{\theta_d} J^{(D)}$$

- **Gradient descent** for the generator

$$J^{(G)} = -\frac{1}{2} \mathbb{E}_{z \sim p_{model}} \log \left(1 - D_{\theta_d}(G_{\theta_g}(z)) \right)$$

$$\theta_g \leftarrow \arg \min_{\theta_g} J^{(G)}$$



Optimization in NN

Take **k** gradient steps for the discriminator (k a hyperparameter), each doing the following:

- Sample m noise samples, $\{z^{(1)}, z^{(2)}, \dots, z^{(m)}\}$ from $p_{\text{model}}(z)$.
- Sample m actual samples, $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ from $p_{\text{data}}(x)$: (a minibatch of your input data.)
- Perform an optimization step on the **discriminator**:



Optimization in NN

Take **k** gradient steps for the discriminator (k a hyperparameter), each doing the following:

- Sample m noise samples, $\{z^{(1)}, z^{(2)}, \dots, z^{(m)}\}$ from $p_{\text{model}}(z)$.
- Sample m actual samples, $\{x^{(1)}, x^{(2)}, \dots, x^{(m)}\}$ from $p_{\text{data}}(x)$: (a minibatch of your input data.)
- Perform an optimization step on the **discriminator**:

Do gradient descent step for the generator:

- Sample m noise samples, $\{z^{(1)}, z^{(2)}, \dots, z^{(m)}\}$ from $p_{\text{model}}(z)$.
- Perform an optimization step on the **generator**:



Optimization in NN

For epoch in range(epochs):

For batch in batches:

Compute $D_{\theta_d}(x)$ vs Real labels

Compute $D_{\theta_d}(G_{\theta_g}(z))$ vs Fake labels

Compute $J^{(D)}$

Backpropagation

Update network

Discriminator Network

For epoch in range(epochs):

For batch in batches:

Compute $D_{\theta_d}(G_{\theta_g}(z))$ vs Real labels

Compute $J^{(G)}$

Backpropagation

Update network

Generator Network




GAN Applications: Original GAN



Goodfellow et al. (2014) Generative Adversarial Nets



DCGAN



Radford, Alec, Luke Metz, and Soumith Chintala. "Unsupervised representation learning with deep convolutional generative adversarial networks." (2015).



Similarities in Hidden Space



woman with glasses



Text to Image Synthesis

this small bird has a pink breast and crown, and black primaries and secondaries.



this magnificent fellow is almost all black with a red crest, and white cheek patch.



the flower has petals that are bright pinkish purple with white stigma



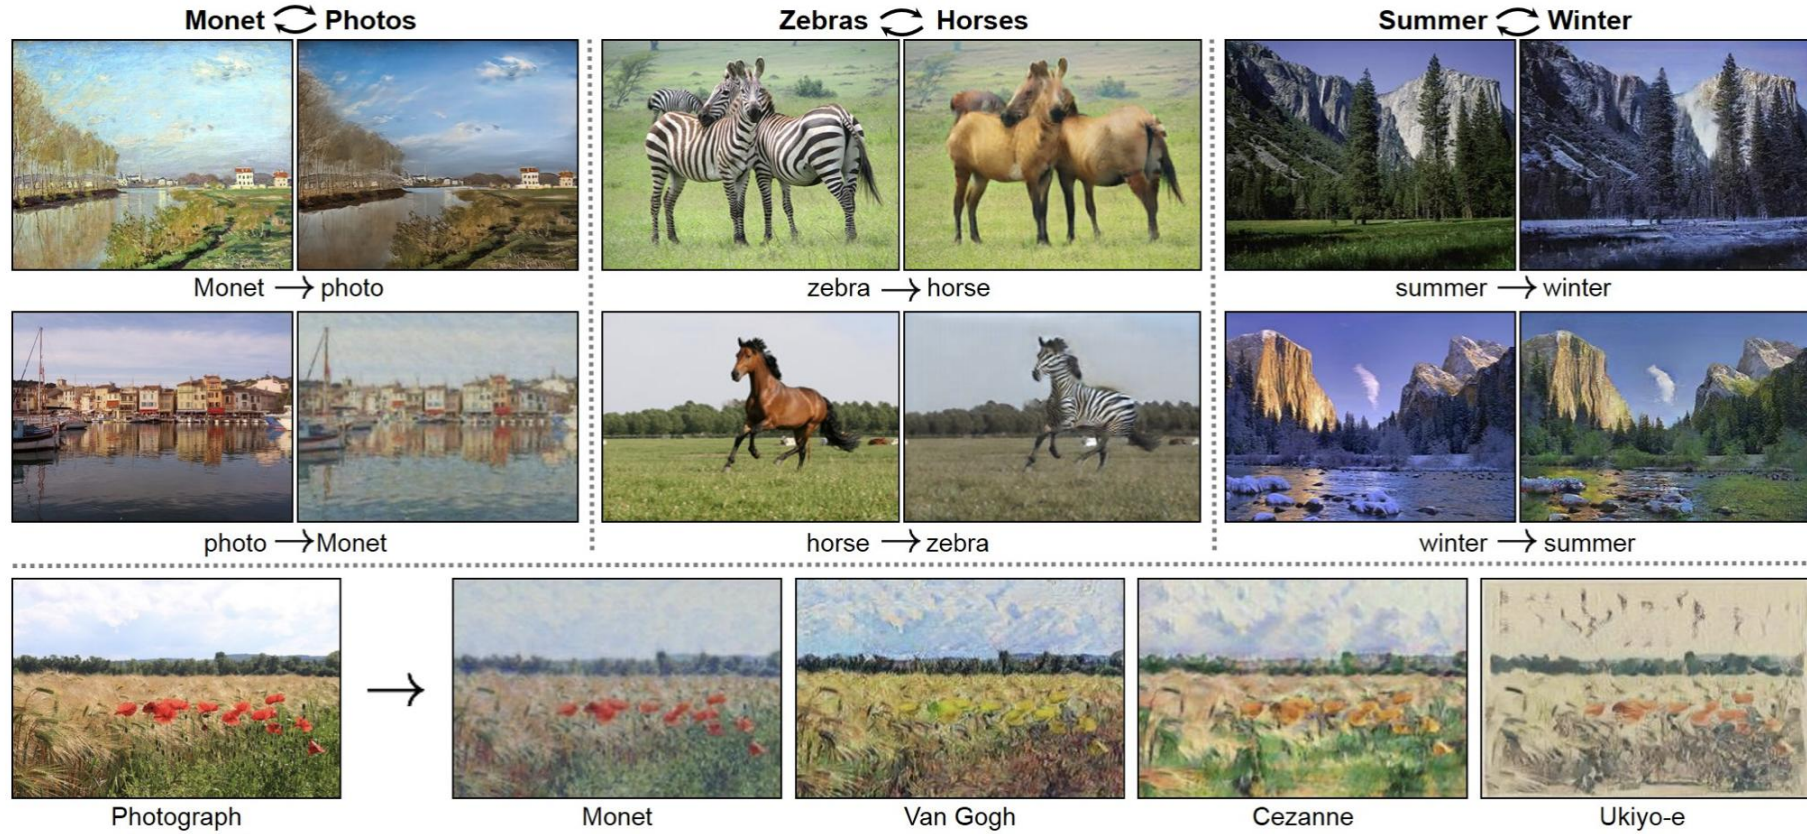
this white and yellow flower have thin white petals and a round yellow stamen



Reed et al. Generative Adversarial Text to Image Synthesis (2017)



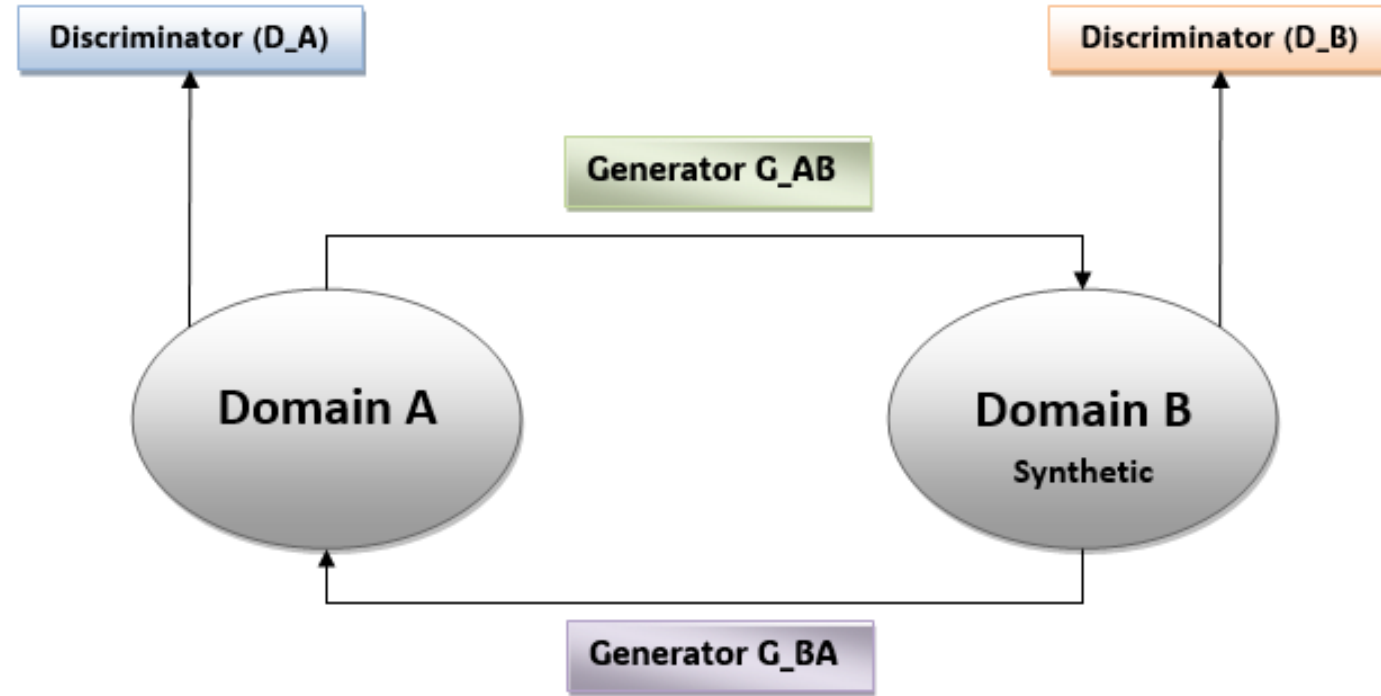
CycleGAN



Zhu et al., Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, ICCV 2017

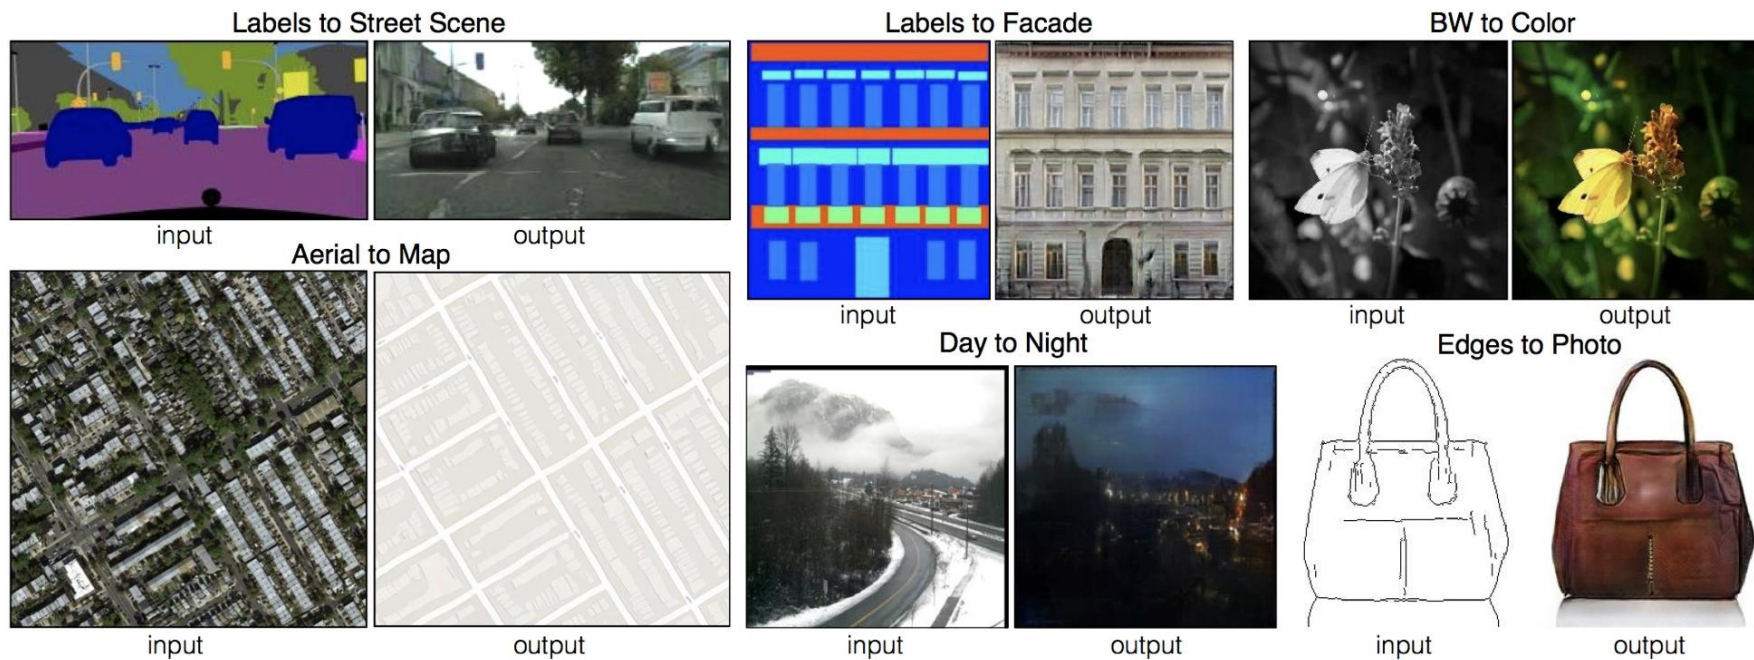


CycleGAN





Pix2Pix



P. Isola et al. Image-to-Image Translation with Conditional Adversarial Nets, CVPR 2017




GAN Example:

MNIST Generation with Vanilla GAN



MNIST Generation with GAN



**Before
Training**



**After
Training**



Prepare Data

```
from torchvision.datasets import MNIST
from torch.utils.data import DataLoader
from torchvision import transforms

# Define a transformation to convert the data into Tensors
train_transforms = transforms.Compose([transforms.ToTensor()])

# Download the train and test MNIST data and transform it into Tensors
train_data = MNIST(root="./train.", train=True, download=True, transform=train_transforms)
```

Load MNIST dataset and in built
DataLoader from torch

Configure the data transformation (to
PyTorch tensors)

Download the training and testing data
N = 60000



Define Model (Generator)

```
class Generator(torch.nn.Module):

    def __init__(self, batchsize, input_noise_dim):

        super(Generator, self).__init__()

        self.batchsize = batchsize # Batch size for input data
        self.input_noise_dim = input_noise_dim # Dimension of the input data

        self.fc1 = torch.nn.Linear(input_noise_dim, 128) # Fully connected Layer 1
        self.LReLU = torch.nn.LeakyReLU() # Leaky ReLU activation function
        self.fc2 = torch.nn.Linear(128, 1 * 28 * 28) # Fully connected Layer 2
        self.output = torch.nn.Tanh() # Hyperbolic Tangent activation function

    def forward(self, x):

        layer1 = self.LReLU(self.fc1(x)) # Apply Leaky ReLU to the first fully connected layer
        layer2 = self.output(self.fc2(layer1)) # Apply Tanh to the second fully connected layer
        out = layer2.view(self.batchsize, 1, 28, 28) # Reshape the output to match image dimensions

        return out
```

Takes batchsize and input noise dimension as inputs

Define FC1, FC2 layers
Uses LeakyReLU() as hidden layer activation
Uses Tanh() as output layer activation

Define signal propagation
input noise -> FC1 -> FC2 -> Output



Define Model (Discriminator)

```
class Discriminator(torch.nn.Module):

    def __init__(self, batchsize):

        super(Discriminator, self).__init__()

        self.batchsize = batchsize # Batch size for input data

        self.fc1 = torch.nn.Linear(1 * 28 * 28, 128) # Fully connected layer 1
        self.LReLU = torch.nn.LeakyReLU() # Leaky ReLU activation function
        self.fc2 = torch.nn.Linear(128, 1) # Fully connected layer 2
        self.output = torch.nn.Sigmoid() # Sigmoid activation function

    # Function for forward propagation
    def forward(self, x):

        flat = x.view(self.batchsize, -1) # Flatten the input image
        layer1 = self.LReLU(self.fc1(flat)) # Apply Leaky ReLU to the first fully connected layer
        out = self.output(self.fc2(layer1)) # Apply Sigmoid to the second fully connected layer

        return out.view(-1, 1).squeeze(1) # Flatten the output and remove unnecessary dimension
```

Takes batchsize as inputs

Define FC1, FC2 layers

Uses LeakyReLU() as hidden layer activation

Uses sigmoid() as output layer activation

Define signal propagation

input image -> FC1 -> FC2 -> (0/1)

Use .squeeze() to reduce output to 0/1



Define Hyperparameters

```
# Fix random seed
torch.manual_seed(55)

# Define Learning rate + epochs
learning_rate = 0.001
epochs = 5

# Define batch size and num_features/timestep (this is simply the last dimension of train_output_seqs)
batchsize = 128
input_noise_dim = 100

# Create a Discriminator model
disc = Discriminator(batchsize)
gen = Generator(batchsize, input_noise_dim)

# Binary Cross Entropy (BCE) Loss function
loss_func = torch.nn.BCELoss()
optimizer_disc = torch.optim.Adam(disc.parameters(), lr=learning_rate, weight_decay=1e-05)
optimizer_gen = torch.optim.Adam(gen.parameters(), lr=learning_rate, weight_decay=1e-05)

# Determine the device for training (GPU if available, otherwise CPU)
device = torch.device("cuda") if torch.cuda.is_available() else torch.device("cpu")
disc.to(device)
gen.to(device)
```

Define learning rate and epoch number

Define batchsize and input noise dimensions

Define Discriminator and Generator networks

Using Binary Cross-Entropy loss and Adam Optimizer with L2 regularization

Device for training (GPU/CPU)



Identify Tracked Values

```
gen_train_loss_list = []  
disc_train_loss_list = []
```

Lists for storing
generator/discriminator training loss



Train Model

```
# Create DataLoader objects to efficiently load the training and test data in batches  
train_loader = DataLoader(train_data, batch_size=batchsize, shuffle=False, drop_last=True)
```

Define the data loader for training and testing

```
# Set the device as CUDA or CPU based on availability  
if torch.cuda.is_available():  
    device = torch.device("cuda")  
else:  
    torch.device("cpu")
```

Define device for training

```
# Run training for each epoch
```

```
for epoch in range(epochs):
```

```
    print('Epoch {}/{}'.format(epoch + 1, epochs))  
    running_loss_D = 0  
    running_loss_G = 0
```

```
    for inputs, labels in train_loader:
```

```
        inputs = inputs.to(device)
```

```
        # Convert labels into torch tensors with the proper size as per the batch size  
        real_label = torch.full((batchsize,), 1, dtype=inputs.dtype, device=device)  
        fake_label = torch.full((batchsize,), 0, dtype=inputs.dtype, device=device)
```

Initialize Discriminator/Generator loss for given epoch

Create labels for real (1) vs fake (0) data



Train Model (Discriminator)

```
# Zero the gradients of the Discriminator optimizer
optimizer_disc.zero_grad()

# Compute output from the Discriminator
output = disc(inputs)

# Discriminator real loss
D_real_loss = loss_func(output, real_label)
D_real_loss.backward()

# Generate random noise data as input to the Generator
noise = torch.randn(batchsize, input_noise_dim, device=device)

# Generate fake images using the Generator
fake = gen(noise)

# Pass fake images through the Discriminator with gradient detachment
output = disc(fake.detach())

# Discriminator fake loss
D_fake_loss = loss_func(output, fake_label)
D_fake_loss.backward()

# Total Loss for the Discriminator
Disc_loss = D_real_loss + D_fake_loss
running_loss_D += Disc_loss

# Update Discriminator's parameters
optimizer_disc.step()
```

Clear the gradient for discriminator network

Compute the discriminator outputs for given real data inputs

Compute the loss with respect to real labels (1)

Perform back propagation

Initialize noise to be fed into generator network

Compute the fake images from the generator

Feed the generated images to discriminator and get discriminator outputs

Compute the loss with respect to fake labels (0)

Perform back propagation

Final discriminator loss is **loss from real labels + fake labels**

Update discriminator network



Train Model (Generator)

```
# Zero the gradients of the Generator optimizer  
optimizer_gen.zero_grad()
```

Clear the gradient for generator network

```
# Pass fake images obtained from the Generator to the Discriminator  
output = disc(fake)
```

Feed the generated image to discriminator

```
# Calculate Generator loss by giving fake images as input but providing real labels  
Gen_loss = loss_func(output, real_label)  
running_loss_G += Gen_loss
```

Compute complementary discriminator loss with respect to **real labels**

```
# Backpropagation for the Generator  
Gen_loss.backward()
```

Perform back-propagation

```
# Update Generator's parameters  
optimizer_gen.step()
```

Update generator network

```
disc_train_loss_list.append(Disc_loss.item())  
gen_train_loss_list.append(Gen_loss.item())
```

Store the loss the respective lists



Visualize & Evaluate Model

```
import matplotlib.pyplot as plt
```

```
# Function to plot an image
```

```
def show_image(img):
```

```
# Convert the image from a tensor to a NumPy array
```

```
    npimg = img.numpy()
```

```
# Transpose the NumPy array to the correct format for displaying
```

```
    plt.imshow(np.transpose(npimg, (1, 2, 0)))
```

Define the function for converting torch image array to numpy array for plotting

```
import torchvision
```

```
# Generate random noise for generating fake images
```

```
random_noise = torch.randn(128, input_noise_dim, device=device)
```

```
# Generate fake images from the random noise using the Generator
```

```
fake = gen(random_noise)
```

```
fake = fake.cpu() # Move the generated fake images to the CPU for displaying
```

```
# Create a Matplotlib figure and axis for displaying the fake images
```

```
fig, ax = plt.subplots(figsize=(20, 8.5))
```

```
# Display the fake images in a grid (e.g., 10x5 grid)
```

```
show_image(torchvision.utils.make_grid(fake[0:50], 10, 5))
```

```
plt.show()
```


Define the random noise to be fed to generator for testing purpose

Feed the noise the generator to produce outputs and move them to cpu

Plot the first 50 generated images




Visualize & Evaluate Model





Visualize & Evaluate Model



Generated samples




Generator Extension with Convolution:

2D Transpose Convolution




DCGAN Architecture





DCGAN Architecture




Layer (type)	Output Shape	Param #
ConvTranspose2d-1	[-1, 512, 4, 4]	819,200
BatchNorm2d-2	[-1, 512, 4, 4]	1,024
ReLU-3	[-1, 512, 4, 4]	0
ConvTranspose2d-4	[-1, 256, 8, 8]	2,097,152
BatchNorm2d-5	[-1, 256, 8, 8]	512
ReLU-6	[-1, 256, 8, 8]	0
ConvTranspose2d-7	[-1, 128, 16, 16]	524,288
BatchNorm2d-8	[-1, 128, 16, 16]	256
ReLU-9	[-1, 128, 16, 16]	0
ConvTranspose2d-10	[-1, 64, 32, 32]	131,072
BatchNorm2d-11	[-1, 64, 32, 32]	128
ReLU-12	[-1, 64, 32, 32]	0
ConvTranspose2d-13	[-1, 3, 64, 64]	3,072
Tanh-14	[-1, 3, 64, 64]	0

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	3,072
LeakyReLU-2	[-1, 64, 32, 32]	0
Conv2d-3	[-1, 128, 16, 16]	131,072
BatchNorm2d-4	[-1, 128, 16, 16]	256
LeakyReLU-5	[-1, 128, 16, 16]	0
Conv2d-6	[-1, 256, 8, 8]	524,288
BatchNorm2d-7	[-1, 256, 8, 8]	512
LeakyReLU-8	[-1, 256, 8, 8]	0
Conv2d-9	[-1, 512, 4, 4]	2,097,152
BatchNorm2d-10	[-1, 512, 4, 4]	1,024
LeakyReLU-11	[-1, 512, 4, 4]	0
Conv2d-12	[-1, 1, 1, 1]	8,192
Sigmoid-13	[-1, 1, 1, 1]	0
Flatten-14	[-1, 1]	0



Conv2D vs ConvTranspose2D




Conv2D()

input image = (5, 5)

Kernel size = (3, 3)

Output image = (3, 3)



ConvTranspose2D()


input image = (2, 2)

Kernel size = (3, 3)

Output image = (4, 4)



Conv2D vs ConvTranspose2D



`torch.nn.ConvTranspose2d(`

<code>in_channels</code>	# of channels of input
<code>out_channels</code>	# of channels of output
<code>kernel_size</code>	Size of the convolving Filter
<code>stride</code>	Stride of the convolution
<code>Padding</code>	Padding added to input

`)`

$$H_{out} = (H_{in} - 1) * stride - 2 * padding + 1 * (kernel_size - 1) + 1$$

$$W_{out} = (W_{in} - 1) * stride - 2 * padding + 1 * (kernel_size - 1) + 1$$




LAB 8 ASSIGNMENT:

MNIST Generation with DCGAN



MNIST Generation with DCGAN



In this exercise, you will use DCGAN architecture to generate **MNIST hand-written images**.

You are free to design architectures for Generator and Discriminator such as # of convolution layers, activation functions, regularization techniques etc.

You are also free to pick your own hyperparameters e.g., total epochs, batch size, learning rate, optimizer etc

Make sure to use **ConvTranspose2D()** for Generator instead of Conv2d().

After training, plot **training loss for both generator and discriminator** as well as the **50 best generated samples** similar to example task. Comment on how their qualities differ from Vanilla-GAN - Are they better quality?